


Network Flows and Bipartite Matching

Rangaballav Pradhan
ITER, SOADU

Flow Network

- A **Flow Network** is a special type of transportation network where some materials (traffic) flows across the nodes and through the edges.
- For example:
 - a road transportation system in which the edges are road segments and the nodes are junctions;
 - a computer network in which the edges are links that can carry data packets and the nodes are switches;
 - a fluid network in which edges are pipes that carry liquid, and the nodes are junctures where pipes are plugged together.

Intuition. Material flowing through a transportation network; material originates at source (s) and is sent to sink (t).




Flow Network

- A **Flow Network** is a special type of transportation network where some materials (traffic) flows across the nodes and through the edges.
- A flow network is associated with several ingredients:
 - capacities on the edges, indicating how much they can carry;
 - source nodes in the graph, which generate traffic;
 - sink or destination nodes in the graph, which can “absorb” traffic as it arrives;
 - the traffic (flow) itself, which is transmitted across the edges.

Assumptions.

- A single source node (s) with no in-degree and single sink (t) node with no out-degree.
- At least one edge incident to each node;
- All capacities are integers.




Flow Network: Definition

A **Flow Network** is a tuple $G = (V, E, s, t, c)$.

- Digraph (V, E) with source $s \in V$ and sink $t \in V$.
- Capacity $c(e) \geq 0$ for each $e \in E$.

assume all nodes are reachable from s




Flow: Definition

Def. An *st-flow* (flow) $f: E \rightarrow R$ is a function that satisfies:

- For each $e \in E$: $0 \leq f(e) \leq c(e)$ [capacity constraint]
- For each $v \in V - \{s, t\}$: $\sum_{e \text{ in to } v} f(e) = \sum_{e \text{ out of } v} f(e)$ [flow conservation constraint]
Inflow (v) Outflow (v)

- the value $f(e)$ represents the amount of flow carried by edge e .
- Inflow (v): the total flow value $f(e)$ over all edges entering node v .
- Outflow (v): the total flow value $f(e)$ over all edges leaving node v .

inflow at $v = 5 + 5 + 0 = 10$
outflow at $v = 10 + 0 = 10$




Value of Flow

Def. An *st*-flow (flow) f is a function that satisfies:

- For each $e \in E$: $0 \leq f(e) \leq c(e)$ [capacity]
- For each $v \in V - \{s, t\}$: $\sum_{e \text{ in to } v} f(e) = \sum_{e \text{ out of } v} f(e)$ [flow conservation]

Def. The **value** of a flow f is: $val(f) = \sum_{e \text{ out of } s} f(e) - \sum_{e \text{ in to } s} f(e)$




Maximum-Flow Problem

Def. An *st*-flow (flow) f is a function that satisfies:

- For each $e \in E$: $0 \leq f(e) \leq c(e)$ [capacity]
- For each $v \in V - \{s, t\}$: $\sum_{e \text{ in to } v} f(e) = \sum_{e \text{ out of } v} f(e)$ [flow conservation]

Def. The **value** of a flow f is: $val(f) = \sum_{e \text{ out of } s} f(e) - \sum_{e \text{ in to } s} f(e)$

Max-flow problem. Given a Flow network $G = (V, E, s, t, c)$, Find a flow of maximum value.

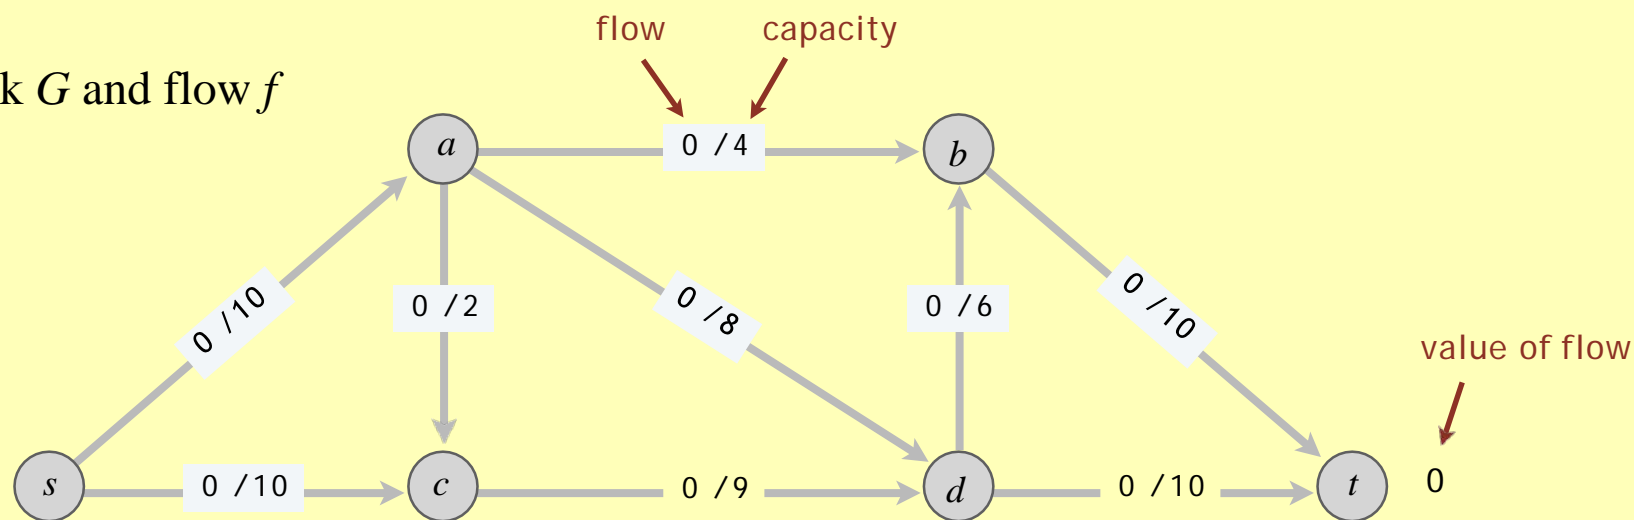


Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

Flow Network G and flow f




Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

Flow Network G and flow f




Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

Flow Network G and flow f




Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

Flow Network G and flow f




Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

Flow Network G and flow f




Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

Ending Flow value = 16

Flow Network G and flow f




Towards a max-flow algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find an $s \rightsquigarrow t$ path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

But maximum flow value = 19

Flow Network G and flow f




Why the greedy algorithm fails?

Q. Why does the greedy algorithm fail?

A. Once greedy algorithm increases flow on an edge, it never decreases it.

Ex. Consider flow network G .

- Greedy algorithm could choose $s \rightarrow u \rightarrow v \rightarrow t$ as first path.




Bottom line. Need some mechanism to “undo” a bad decision.

Residual network

Original edge. $e = (u, v) \in E$.

- Flow $f(e)$.
- Capacity $c(e)$.


Original Flow Network G



Reverse edge. $e^{\text{reverse}} = (v, u)$.

- “Undo” flow sent.

Residual Flow Network G_f



Residual capacity.

$$c_f(e) = \begin{cases} c(e) - f(e) & \text{if } e \in E \\ f(e) & \text{if } e^{\text{reverse}} \in E \end{cases}$$

Residual network. $G_f = (V, E_f, s, t, c_f)$.

$$E_f = \{e : f(e) < c(e)\} \cup \{e^{\text{reverse}} : f(e) > 0\}.$$

edges with positive residual capacity

Augmenting path

Def. An **augmenting path** is a simple $s \rightsquigarrow t$ path in the residual network G_f .

Def. The **bottleneck capacity** of an augmenting path P is the minimum residual capacity of any edge in P .

Key property. Let f be a flow and let P be an augmenting path in G_f . Then, after calling $f' \leftarrow \text{AUGMENT}(f, c, P)$, the resulting f' is a flow and $\text{val}(f') = \text{val}(f) + \text{bottleneck}(G_f, P)$.

AUGMENT(f, c, P)

Let $b = \text{bottleneck}(P, f)$

For each edge $(u, v) \in P$

 If $e = (u, v)$ is a forward edge then

 increase $f(e)$ in G by b

 Else $((u, v)$ is a backward edge, and let $e = (v, u)$)

 decrease $f(e)$ in G by b

 Endif

Endfor

Return(f)

Ford–Fulkerson algorithm

Ford–Fulkerson augmenting path algorithm.

- Start with $f(e) = 0$ for each edge $e \in E$.
- Find a simple $s \rightsquigarrow t$ path P in the residual network G_f .
- Augment flow along path P .
- Repeat until you get stuck.

FORD–FULKERSON(G)

FOREACH edge $e \in E$:

$f(e) \leftarrow 0$.

$G_f \leftarrow$ residual network of G with respect to flow f .


WHILE (there exists an $s \rightsquigarrow t$ path P in G_f)

$f \leftarrow \text{AUGMENT}(f, c, P)$.

Update G_f .


RETURN f .

augmenting path



Time = $O(E \cdot f)$

Example



Ford–Fulkerson algorithm: Analysis

- Let $m = |E|$ and $n = |V|$
- The *for* loop runs for $O(m)$ times.
- The algorithm terminates in at most C iterations of the While loop in the worst case, where $C = \sum_{e \text{ out of } s} c(e)$.
- The residual graph G_f has at most $2m$ edges. Given the new flow f , we can build the new residual graph in $O(m + n)$ time
- We can maintain G_f using an adjacency list representation in $O(m + n)$ time.
- To find an s-t path in G_f , we can use BFS or DFS, which run in $O(m + n)$ time;
- By our earlier assumption that each vertex is incident with at least one edge, i.e., $m \geq n/2$. So, $O(m + n)$ is the same as $O(m)$.
- The procedure $\text{AUGMENT}(f, c, P)$ takes time $O(n)$, as the path P has at most $n - 1$ edges.
- So, the total time complexity in worst case: $O(E * f)$ or $O(m * C)$

FORD–FULKERSON(G)

FOREACH edge $e \in E$:

$f(e) \leftarrow 0$.

$G_f \leftarrow$ residual network of G with respect to flow f .


WHILE (there exists an s-t path P in G_f)

$f \leftarrow \text{AUGMENT}(f, c, P)$.

Update G_f .

RETURN f .

augmenting path




Cut in a Flow Network

Def. An *st-cut (cut)* is a partition (A, B) of the nodes with $s \in A$ and $t \in B$.

Def. Its *capacity* is the sum of the capacities of the edges from A to B .

$$cap(A, B) = \sum_{e \text{ out of } A} c(e)$$




Cut in a Flow Network

Def. An *st-cut (cut)* is a partition (A, B) of the nodes with $s \in A$ and $t \in B$.

Def. Its *capacity* is the sum of the capacities of the edges from A to B .

$$cap(A, B) = \sum_{e \text{ out of } A} c(e)$$




Minimum-cut problem

Def. An *st-cut (cut)* is a partition (A, B) of the nodes with $s \in A$ and $t \in B$.


Def. Its *capacity* is the sum of the capacities of the edges from A to B .

$$cap(A, B) = \sum_{e \text{ out of } A} c(e)$$

Min-cut problem. Find a cut of minimum capacity.



What is the capacity of the cut?



Relationship between flows and cuts

Flow value lemma. Let f be any st -flow and let (A, B) be any st -cut. Then, the value of the st -flow f equals the net flow across the cut (A, B) .

$$val(f) = \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e)$$


Proof. We know, $val(f) = \sum_{e \text{ out of } s} f(e) - \sum_{e \text{ in to } s} f(e)$

Every node v in A is internal except s , and we know that for all such nodes, $f^{out}(v) = f^{in}(v)$

So,

$$val(f) = \sum_{v \in A} (f^{out}(v) - f^{in}(v))$$

$$\sum_{v \in A} f^{out}(v) - f^{in}(v) = \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ into } A} f(e) = f^{out}(A) - f^{in}(A)$$




Relationship between flows and cuts

Flow value lemma. Let f be any st -flow and let (A, B) be any st -cut. Then, the value of the flow f equals the net flow across the cut (A, B) .

$$\text{val}(f) = \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e)$$

net flow across cut = $10 + 5 + 10 = 25$




value of flow = 25

Relationship between flows and cuts

Flow value lemma. Let f be any flow and let (A, B) be any cut. Then, the value of the flow f equals the net flow across the st -cut (A, B) .


$$\text{val}(f) = \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e)$$

$$\text{net flow across cut} = (10 + 10 + 5 + 10 + 0 + 0) - (5 + 5 + 0 + 0) = 25$$



Relationship between flows and cuts

What is the net flow across the given cut?




Relationship between flows and cuts

Weak duality. Let f be any st -flow and (A, B) be any st -cut. Then, $val(f) \leq cap(A, B)$.

Pf.


$$\begin{aligned}
 val(f) &= \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) \\
 &\leq \sum_{e \text{ out of } A} f(e) \\
 &\leq \sum_{e \text{ out of } A} c(e) \\
 &= cap(A, B) \quad \blacksquare
 \end{aligned}$$

flow value lemma



value of flow = 27

\leq



capacity of cut = 30

Relationship between flows and cuts

Theorem 1. If f is an st -flow such that there is no s - t path in the residual graph G_f , then there is an st -cut (A^*, B^*) in G for which $v(f) = c(A^*, B^*)$. Consequently, f has the maximum value of any flow in G , and (A^*, B^*) has the minimum capacity of any st -cut in G .

Pf. The statement claims the existence of the cut (A^*, B^*) satisfying a certain desirable property;

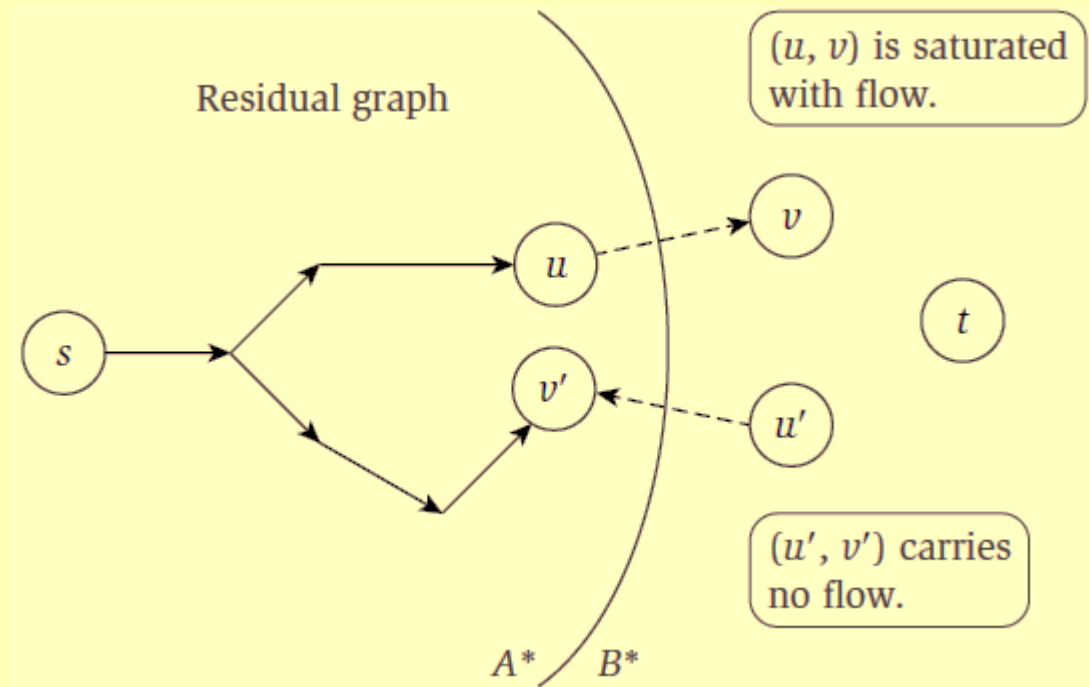
- To prove, we must identify such a cut.
- Let A^* denote the set of all nodes v in G for which there is an s - v path in the final G_f . Let B^* denote the set of all other nodes: $B^* = V - A^*$.
- First we establish that (A^*, B^*) is an s - t cut.
- It is clearly a partition of V . The source s belongs to A^* since there is always a path from s to s .
- Moreover, $t \notin A^*$ by the assumption that there is no s - t path in G_f ; hence $t \in B^*$ as desired.

Relationship between flows and cuts

Theorem 1. If f is an st -flow such that there is no s - t path in the residual graph G_f , then there is an st -cut (A^*, B^*) in G for which $v(f) = c(A^*, B^*)$. Consequently, f has the maximum value of any flow in G , and (A^*, B^*) has the minimum capacity of any st -cut in G .

Pf. Suppose that $e = (u, v)$ is an edge in G for which $u \in A^*$ and $v \in B^*$. We claim that $f(e) = c(e)$.

- For if not, e would be a forward edge in the residual graph G_f , and there will be a path s - u - v in G_f contradicting our assumption that $v \in B^*$.
- So all edges out of A^* are completely saturated with flow.




Relationship between flows and cuts

Theorem 1. If f is an st -flow such that there is no s - t path in the residual graph G_f , then there is an st -cut (A^*, B^*) in G for which $v(f) = c(A^*, B^*)$. Consequently, f has the maximum value of any flow in G , and (A^*, B^*) has the minimum capacity of any st -cut in G .

Pf. Suppose that $e' = (u', v')$ is an edge in G for which $v' \in A^*$ and $u' \in B^*$. We claim that $f(e') = 0$.

- For if not, e' would produce a backward edge $e'' = (v', u')$ in the residual graph G_f , and there will be a path s - v' - u' in G_f contradicting our assumption that $u' \in B^*$.
- All edges into A^* are completely unused.

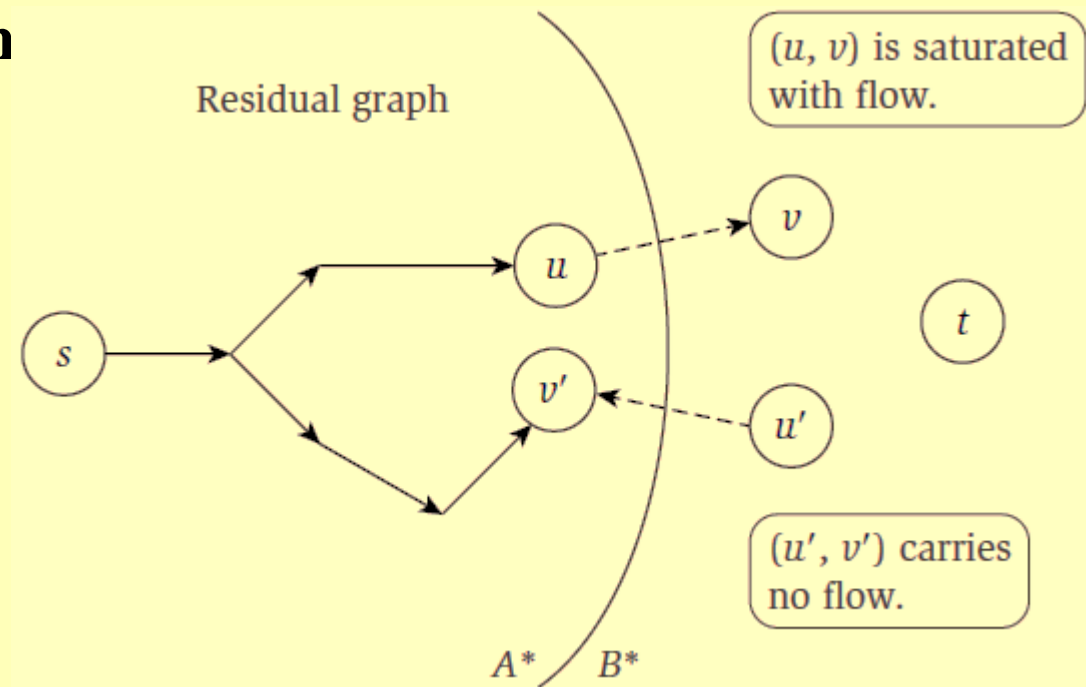


Relationship between flows and cuts

Theorem 1. If f is an st -flow such that there is no s - t path in the residual graph G_f , then there is an st -cut (A^*, B^*) in G for which $v(f) = c(A^*, B^*)$. Consequently, f has the maximum value of any flow in G , and (A^*, B^*) has the minimum capacity of any st -cut in

Pf.

$$\begin{aligned} v(f) &= f^{\text{out}}(A^*) - f^{\text{in}}(A^*) \\ &= \sum_{e \text{ out of } A^*} f(e) - \sum_{e \text{ into } A^*} f(e) \\ &= \sum_{e \text{ out of } A^*} c_e - 0 \\ &= c(A^*, B^*). \quad \blacksquare \end{aligned}$$



Max-flow min-cut theorem

Augmenting path theorem. A flow f is a max flow iff. no augmenting paths.

Max-flow min-cut theorem. In every flow network, the maximum value of an st -flow is equal to the minimum capacity of an st -cut.

Pf. The point is that f in **Theorem 1** must be a maximum st -flow;

- If there were a flow $f' > f$, the value of f' would exceed the capacity of (A, B) , and this would contradict the **weak duality**.
- Similarly, it follows that (A, B) in **Theorem 1** is a minimum cut and no other cut can have smaller capacity.
- If there were a cut (A', B') of smaller capacity, it would be less than the value of f , and this again would contradict the **weak duality**.
- Due to these implications, the Max-Flow Min-Cut Theorem holds its claim.


Computing a minimum cut from a maximum flow

Theorem. Given any max flow f , can compute a min cut (A, B) in $O(m)$ time (m = no. of edges in the residual graph).

Pf. Let A = set of nodes reachable from s in residual network G_f using BFS/DFS.


▪

argument from previous slide implies that capacity of (A, B) = value of flow f




Example

Find the maximum flow and minimum cut capacity in the given graph, identify the cut.



Choosing good augmenting path

- The worst case time complexity of Ford-Fulkerson: $O(m \cdot C)$
- It is a reasonable bound for smaller C value;
- However, it is very weak when C is large.
- For example, if Ford-Fulkerson algorithm selects the augmenting paths in the following order:
 - $s-u-v-t$
 - $s-v-u-t$
 - $s-u-v-t$
 - $s-v-u-t$
 -
- The running time will be significantly high.



Choosing good augmenting path: History

year	method	# augmentations	running time	
1955	augmenting path	$n C$	$O(m n C)$	fat paths
1972	fattest path	$m \log (mC)$	$O(m^2 \log n \log (mC))$	
1972	capacity scaling	$m \log C$	$O(m^2 \log C)$	
1985	improved capacity scaling	$m \log C$	$O(m n \log C)$	
1970	shortest augmenting path	$m n$	$O(m^2 n)$	shortest paths
1970	level graph	$m n$	$O(m n^2)$	
1983	dynamic trees	$m n$	$O(m n \log n)$	

Max-flow algorithms: History

year	method	worst case	discovered by
1951	simplex	$O(m n^2 C)$	Dantzig
1955	augmenting paths	$O(m n C)$	Ford–Fulkerson
1970	shortest augmenting paths	$O(m n^2)$	Edmonds–Karp, Dinitz
1974	blocking flows	$O(n^3)$	Karzanov
1983	dynamic trees	$O(m n \log n)$	Sleator–Tarjan
1985	improved capacity scaling	$O(m n \log C)$	Gabow
1988	push-relabel	$O(m n \log (n^2 / m))$	Goldberg–Tarjan
1998	binary blocking flows	$O(m^{3/2} \log (n^2 / m) \log C)$	Goldberg–Rao
2013	compact networks	$O(m n)$	Orlin
2014	interior–point methods	$\tilde{O}(m n^{1/2} \log C)$	Lee–Sidford
2016	electrical flows	$\tilde{O}(m^{10/7} C^{1/7})$	Mądry
20xx		???	

max-flow algorithms with m edges, n nodes, and integer capacities between 1 and C

Choosing good augmenting path

- A natural idea is to select the path with largest bottleneck capacity.
- However, to find such paths can slow down each individual iteration by quite a bit.
- We will avoid this slowdown by not worrying about selecting the path that has exactly the largest bottleneck capacity. Instead, we will maintain a scaling parameter Δ , and we will look for paths that have bottleneck capacity of at least Δ .
- Let $G_f(\Delta)$ be the subgraph of the residual graph G_f consisting only of edges with residual capacity of at least Δ .
- We will work with values of Δ that are powers of 2 and not greater than C .

Scaling-Max-Flow Algorithm

Scaling-Max-Flow(G, s, t, c)

FOREACH edge $e \in E : f(e) \leftarrow 0$.

$\Delta \leftarrow$ largest power of 2 $\leq C$.

WHILE ($\Delta \geq 1$)

$G_f(\Delta) \leftarrow \Delta$ -residual network of G with respect to flow f .

WHILE (there exists an $s \rightsquigarrow t$ path P in $G_f(\Delta)$)

$f \leftarrow \text{AUGMENT}(f, c, P)$.

Update $G_f(\Delta)$.

$\Delta \leftarrow \Delta / 2$.

Δ -scaling phase

RETURN f .

Scaling-Max-Flow Algorithm: Analysis

- **Assumption.** All edge capacities are integers between 1 and C .
- **Lemma 1.** There are $1 + \lfloor \log C \rfloor$ scaling phases.
- **Pf.** Initially $C / 2 < \Delta \leq C$; Δ decreases by a factor of 2 in each iteration. So, total iterations will be $1 + \lfloor \log C \rfloor$.
- **Lemma 2.** Let f be the flow at the end of a Δ -scaling phase. Then, the max-flow value $\leq \text{val}(f) + m \Delta$.

Pf. At the end of a Δ -scaling phase, let's identify a cut (A, B) where A denotes the set of all nodes v in G for which there is an s - v path in $G_f(\Delta)$ and B denotes the set of all other nodes: $B = V - A$.


- We can observe that (A, B) is indeed an st-cut as otherwise the phase would not have ended.

Scaling-Max-Flow Algorithm: Analysis

- **Lemma 2.** Let f be the flow at the end of a Δ -scaling phase. Then, the max-flow value $\leq \text{val}(f) + m \Delta$.

Pf. We can easily realize that all edges e out of A are almost saturated and they satisfy $c_e < f(e) + \Delta$ and all edges into A are almost empty and they satisfy $f(e) < \Delta$.

$$\begin{aligned} v(f) &= \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ into } A} f(e) \\ &\geq \sum_{e \text{ out of } A} (c_e - \Delta) - \sum_{e \text{ into } A} \Delta \\ &= \sum_{e \text{ out of } A} c_e - \sum_{e \text{ out of } A} \Delta - \sum_{e \text{ into } A} \Delta \\ &\geq c(A, B) - m\Delta. \end{aligned}$$



So, $c(A, B) \leq v(f) + m\Delta \Rightarrow \text{max-flow} \leq v(f) + m\Delta$

Scaling-Max-Flow Algorithm: Analysis

- **Lemma 3.** The number of augmentations in a scaling phase is at most $2m$.

Pf. The statement is clearly true in the first scaling phase: we can use each of the edges out of s only for at most one augmentation in that phase.

- Now consider a later scaling phase Δ , and let f_p be the flow at the end of the previous scaling phase. In that phase, we used $\Delta' = 2\Delta$ as our scaling parameter.
- By **Lemma 2**, the maximum flow f^* is at most $v(f^*) \leq v(f_p) + m \Delta' = v(f_p) + 2m\Delta$.
- In the Δ -scaling phase, each augmentation increases the flow by at least Δ , and hence there can be at most $2m$ augmentations..

Scaling-Max-Flow Algorithm: Analysis

- **Lemma 4.** The Scaling Max-Flow Algorithm in a graph with m edges and integer capacities finds a maximum flow in at most $2m(1 + \lceil \log C \rceil)$ augmentations. It can be implemented to run in at most $O(m^2 \log C)$ time.

Pf. Lemma 1 + Lemma 3 $\Rightarrow O(m \log C)$ augmentations.

- Finding an augmenting path takes $O(m)$ time.
- So, the running time of the Scaling Max-Flow Algorithm is $O(m^2 \log C)$.

Scaling-Max-Flow(G, s, t, c)

FOREACH edge $e \in E : f(e) \leftarrow 0$.

$\Delta \leftarrow$ largest power of 2 $\leq C$.

WHILE ($\Delta \geq 1$)

$G_f(\Delta) \leftarrow \Delta$ -residual network of G with respect to flow f .

WHILE (there exists an $s \rightsquigarrow t$ path P in $G_f(\Delta)$)

$f \leftarrow \text{AUGMENT}(f, c, P)$.

Update $G_f(\Delta)$.

$\Delta \leftarrow \Delta / 2$.

Δ -scaling phase

RETURN f .

Max-flow and min-cut applications


Max-flow and min-cut problems are widely applicable model.

- Data mining.
- Open-pit mining.
- Bipartite matching.
- Network reliability.
- Baseball elimination.
- Image segmentation.
- Network connectivity.
- Markov random fields.
- Distributed computing.
- Security of statistical data.
- Egalitarian stable matching.
- Network intrusion detection.
- Multi-camera scene reconstruction.
- Sensor placement for homeland security.
- Many, many, more.

Matching

Def. Given an undirected graph $G = (V, E)$, subset of edges $M \subseteq E$ is a **matching** if each node appears in at most one edge in M .


Max matching. Given a graph G , find a max-cardinality matching.



Bipartite matching

Def. A graph G is **bipartite** if the nodes can be partitioned into two subsets L and R such that every edge connects a node in L with a node in R .


Bipartite matching. Given a bipartite graph $G = (L \cup R, E)$, find a max-cardinality matching.



Bipartite matching: max-flow formulation

Formulation.

- Create digraph $G' = (L \cup R \cup \{s, t\}, E')$ from the given bipartite graph $G = (L \cup R, E)$.
- Direct all edges from L to R , and assign the capacity as 1 for each.
- Add unit-capacity edges from s to each node in L .
- Add unit-capacity edges from each node in R to t .
- Now, run the Ford-Fulkerson Algorithm and find the max flow value which gives the maximum cardinality of bipartite matching.



Max-flow formulation: proof of correctness



- Suppose there is a matching in G consisting of k edges of the form (x, y) where $x \in L$ and $y \in R$.
- Consider the flow f that sends one unit along each path of the form $s-x-y-t$, i.e., $f(e) = 1$ for each edge on these paths. We can verify easily that the capacity and conservation conditions are indeed met and that f is an $s-t$ flow of value k .
- Conversely, suppose there is a flow f' in G' of value k . We know there is an integer-valued flow f of value k ; and since all capacities are 1, this means that $f(e)$ is equal to either 0 or 1 for each edge e .
- Now, consider the set M' of edges of the form (x, y) on which the flow value is 1.

Max-flow formulation: proof of correctness

Observation 1. M' contains k edges where k is the max flow.

Pf. To prove this, consider the cut (A, B) in G' with $A = \{s\} \cup L$.

- The value of the flow is the total flow leaving A , minus the total flow entering A .
- The first of these terms is simply the cardinality of M' since these are the edges leaving A that carry flow, and each carries exactly one unit of flow.
- The second of these terms is 0, since there are no edges entering A .
- Thus, M' contains k edges.





Max-flow formulation: proof of correctness

Observation 2. Each node in L is the tail of at most one edge in M' .

Pf. To prove this, suppose $x \in L$ were the tail of at least two edges in M' .

- Since our flow is integer-valued, this means that at least two units of flow leave from x .
- By conservation of flow, at least two units of flow would have to come into x —but this is not possible, since only a single edge of capacity 1 enters x .
- Thus x is the tail of at most one edge in M' .




Max-flow formulation: proof of correctness

Observation 3. Each node in R is the head of at most one edge in M' .

Pf. By the same reasoning as previous.

Conclusion. The size of the maximum matching in G is equal to the value of the maximum flow in G' . and the edges in such a matching in G are the edges that carry flow from L to R in G' .



Network flow II: quiz

What is running time of Ford–Fulkerson algorithms to find a max- cardinality matching in a bipartite graph with $|L| = |R| = n$?

- A. $O(m + n)$
- B. $O(mn)$
- C. $O(mn^2)$
- D. $O(m^2n)$

Perfect matchings in bipartite graphs


Def. Given a graph $G = (V, E)$, a subset of edges $M \subseteq E$ is a **perfect matching** if each node of G appears in exactly one edge in M .

- Condition for perfect matching: $|L| = |R| = |M|$

Q. When does a bipartite graph have a perfect matching?

Structure of bipartite graphs with perfect matchings.

- Clearly, we must have $|L| = |R|$.



Perfect matchings in bipartite graphs


Notation. Let S be a subset of nodes, and let $N(S)$ be the set of nodes adjacent to nodes in S .

Observation. If a bipartite graph $G = (L \cup R, E)$ has a perfect matching, then $|N(S)| \geq |S|$ for all subsets $S \subseteq L$.

Pf. Each node in S has to be matched to a different node in $N(S)$. ■

$$S = \{2, 4, 5\}$$

$$N(S) = \{2', 5'\}$$



no perfect matching